

25 YEARS ANNIVERSARY
SOICT

HA NOI UNIVERSITY OF SCIENCE AND TECHNOLOGY
SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY



HA NOI UNIVERSITY OF SCIENCE AND TECHNOLOGY
SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

Chapter 7

Stream processing


Structured streaming in Spark

Pain points with DStreams

- Processing with event-time, dealing with late data
 - DStream API exposes batch time, hard to incorporate event-time
- Interoperate streaming with batch AND interactive
 - RDD/DStream has similar API, but still requires translation
- Reasoning about end-to-end guarantees
 - Requires carefully constructing sinks that handle failures correctly
 - Data consistency in the storage while being updated


New model

- Input: data from source as an append-only table
- Trigger: how frequently to check input for new data
- Query: operations on input usual map/filter/reduce new window, session ops



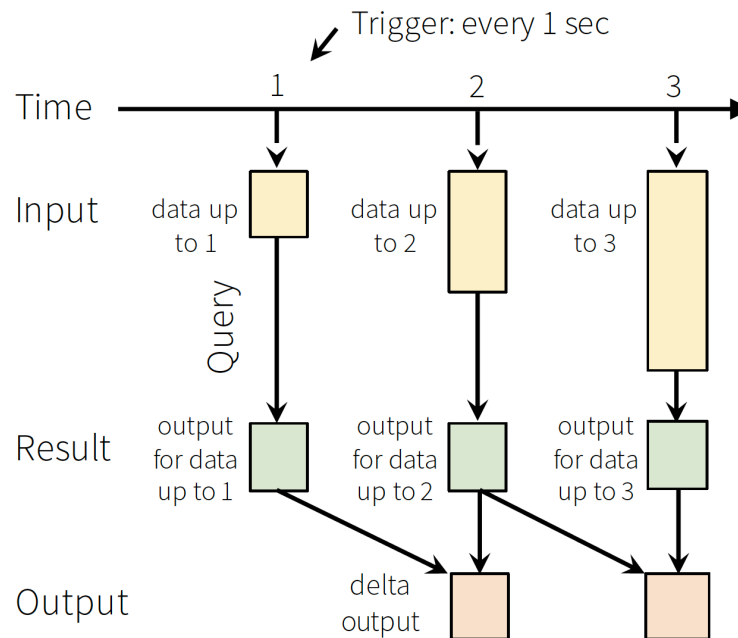
New model (2)

- Result: final operated table updated every trigger interval
- Output: what part of result to write to data sink after every trigger
- Complete output: Write full result table every time



New model (3)

- Delta output: Write only the rows that changed in result from previous batch
- Append output: Write only new rows
- *Not all output modes are feasible with all queries



Batch ETL with DataFrames

```
input = spark.read
    .format("json")
    .load("source-path")
result = input
    .select("device",
            "signal")
    .where("signal > 15")
result.write
    .format("parquet")
    .save("dest-path")
```

- Read from Json file
- Select some devices
- Write to parquet file

Streaming ETL with DataFrames

```
input = spark.read
    .format("json")
    .stream("source-
path")
result = input
    .select("device",
"signal")
    .where("signal >
15")
result.write
    .format("parquet")
    .startStream("dest-
path")
```

- Read from Json file stream
 - Replace load() with stream()
- Select some devices
 - Code does not change
- Write to Parquet file stream
 - Replace save() with startStream()

Streaming ETL with DataFrames

```
input = spark.read  
    .format("json")  
    .stream("source-  
path")
```


```
result = input  
    .select("device",  
"signal")  
    .where("signal >  
15")
```

```
result.write  
    .format("parquet")  
    .startStream("dest-  
path")
```

- `read...stream()` creates a streaming DataFrame, does not start any of the computation
- `write...startStream()` defines where & how to output the data and starts the processing

Streaming ETL with DataFrames

```
input = spark.read
    .format("json")
    .stream("source-path")
result = input
    .select("device",
            "signal")
    .where("signal >
15")
result.write
    .format("parquet")
    .startStream("dest-path")
```



Continuous Aggregations

```
input.avg("signal")
```

- Continuously compute average signal across all devices

```
input.groupBy("device-  
type")  
    .avg("signal")
```

- Continuously compute average signal of each type of device

Continuous Windowed Aggregations

```
input.groupBy(  
    $"device-type",  
    window($"event-  
time-col", "10 min"))  
    .avg("signal")
```

- Continuously compute average signal of each type of device in last 10 minutes using event-time
- Simplifies event-time stream processing (not possible in DStreams)
Works on both, streaming and batch jobs

Joining streams with static data

```
kafkaDataset = spark.read  
    .kafka("iot-updates")  
    .stream()
```

- Join streaming data from Kafka with static data via JDBC to enrich the streaming data ...

```
staticDataset = ctx.read  
    .jdbc("jdbc://", "iot-  
device-info")
```

```
joinedDataset =  
    kafkaDataset.join(  
        staticDataset,  
        "device-type")
```

- ... without having to think that you are joining streaming data

Output modes

Defines what is outputted every time there is a trigger
Different output modes make sense for different queries

- Append mode with non-aggregation queries

```
input.select("device", "signal")  
    .write  
    .outputMode("append")  
    .format("parquet")  
    .startStream("dest-path")
```

- Complete mode with aggregation queries

```
input.agg(count("*"))  
    .write  
    .outputMode("complete"  
)  
    .format("parquet")  
    .startStream("dest-path")
```

Query Management

```
query = result.write  
    .format("parquet")  
    .outputMode("append"  
)  
    .startStream("dest-  
path")
```

```
query.stop()  
query.awaitTermination()  
query.exception()
```

```
query.sourceStatuses()  
query.sinkStatus()
```

- query: a handle to the running streaming computation for managing it
 - Stop it, wait for it to terminate
 - Get status
 - Get error, if terminated
- Multiple queries can be active at the same time
- Each query has unique name for keeping track

Query execution

- Logically
 - Dataset operations on table (i.e. as easy to understand as batch)
- Physically
 - Spark automatically runs the query in streaming fashion (i.e. incrementally and continuously)

