

25 YEARS ANNIVERSARY  
SOICT

HA NOI UNIVERSITY OF SCIENCE AND TECHNOLOGY  
SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY



HA NOI UNIVERSITY OF SCIENCE AND TECHNOLOGY  
SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

# Chapter 4

# NoSQL - part 1

## CAP theorem

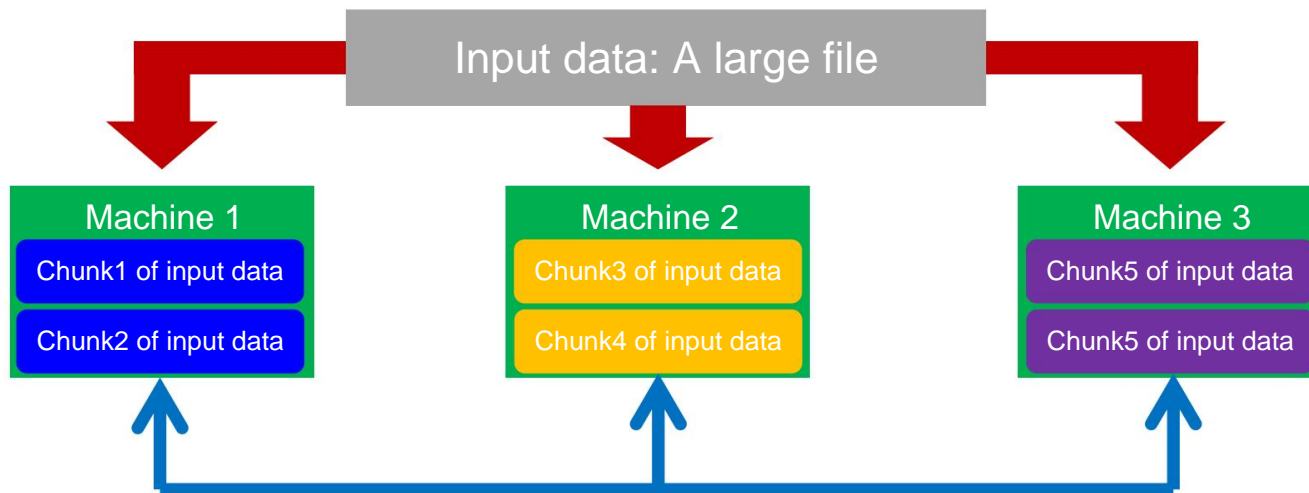
# Scaling Traditional Databases

- Traditional RDBMSs can be either scaled:
  - Vertically (or Up)
    - Can be achieved by hardware upgrades (eg, faster CPU, more memory, or larger disk)
    - Limited by the amount of CPU, RAM and disk that can be configured on a single machine
  - Horizontally (or Out)
    - Can be achieved by adding more machines
    - Requires database sharding and probably replication
    - Limited by the Read-to-Write ratio and overhead communication

# Data sharding

- Data is typically sharded (or striped) to allow for concurrent/parallel accesses •

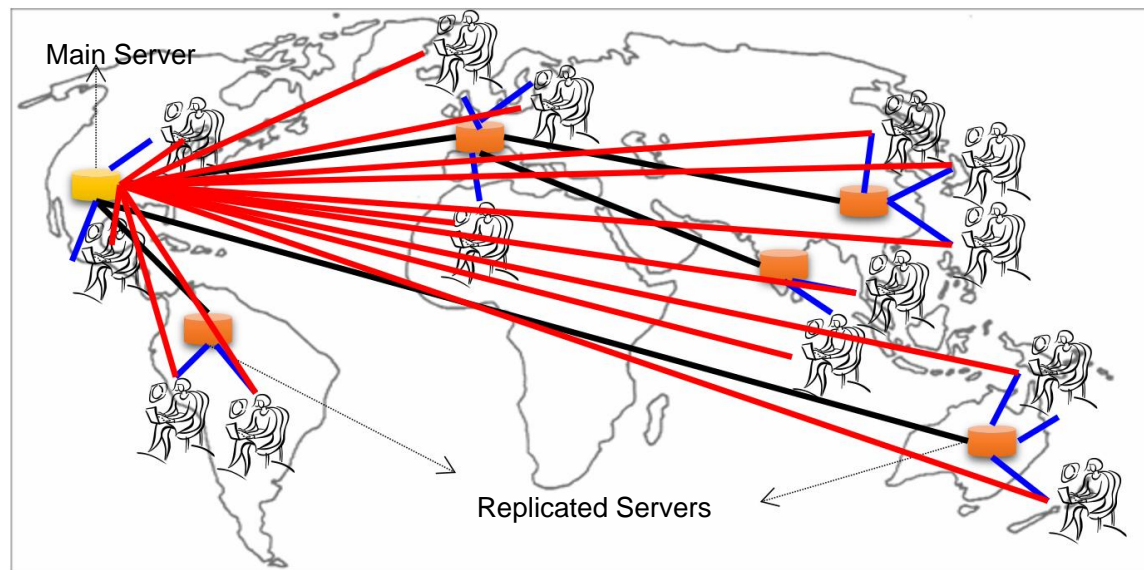
Will it scale for complex query processing?



Eg, Chunks 1, 3 and 5 can be accessed in parallel

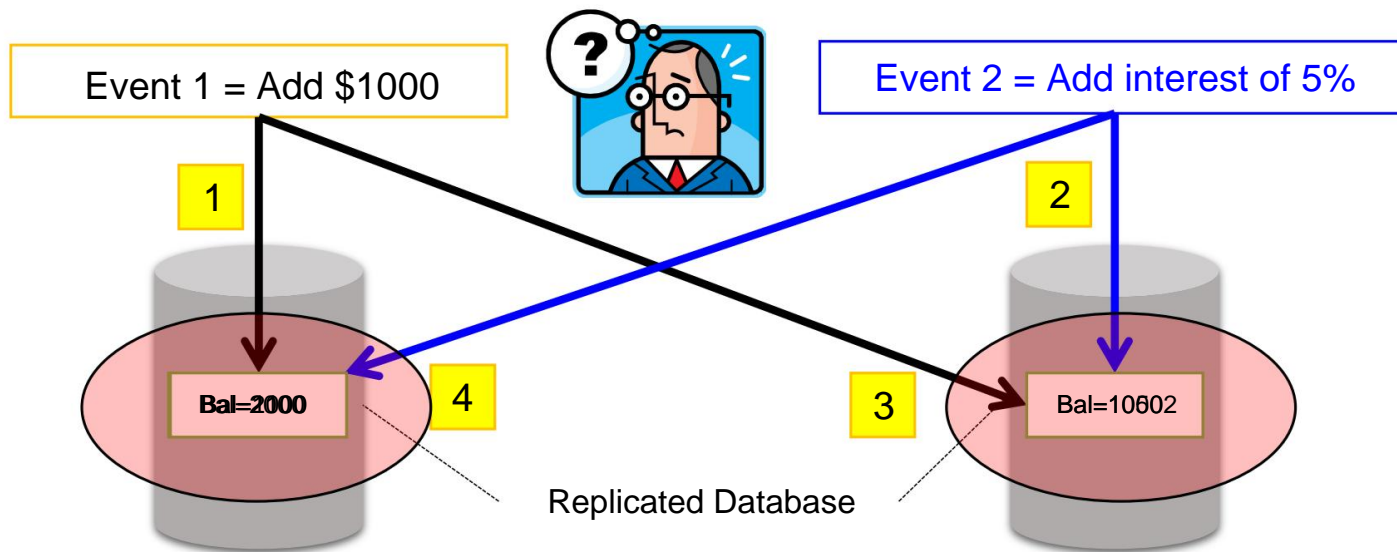
# Data replication

- Replicating data across servers helps in:
  - Avoiding performance bottlenecks
  - Avoiding single point of failures
  - And, therefore, enhancing scalability and availability



# But, Consistency Becomes a Challenge

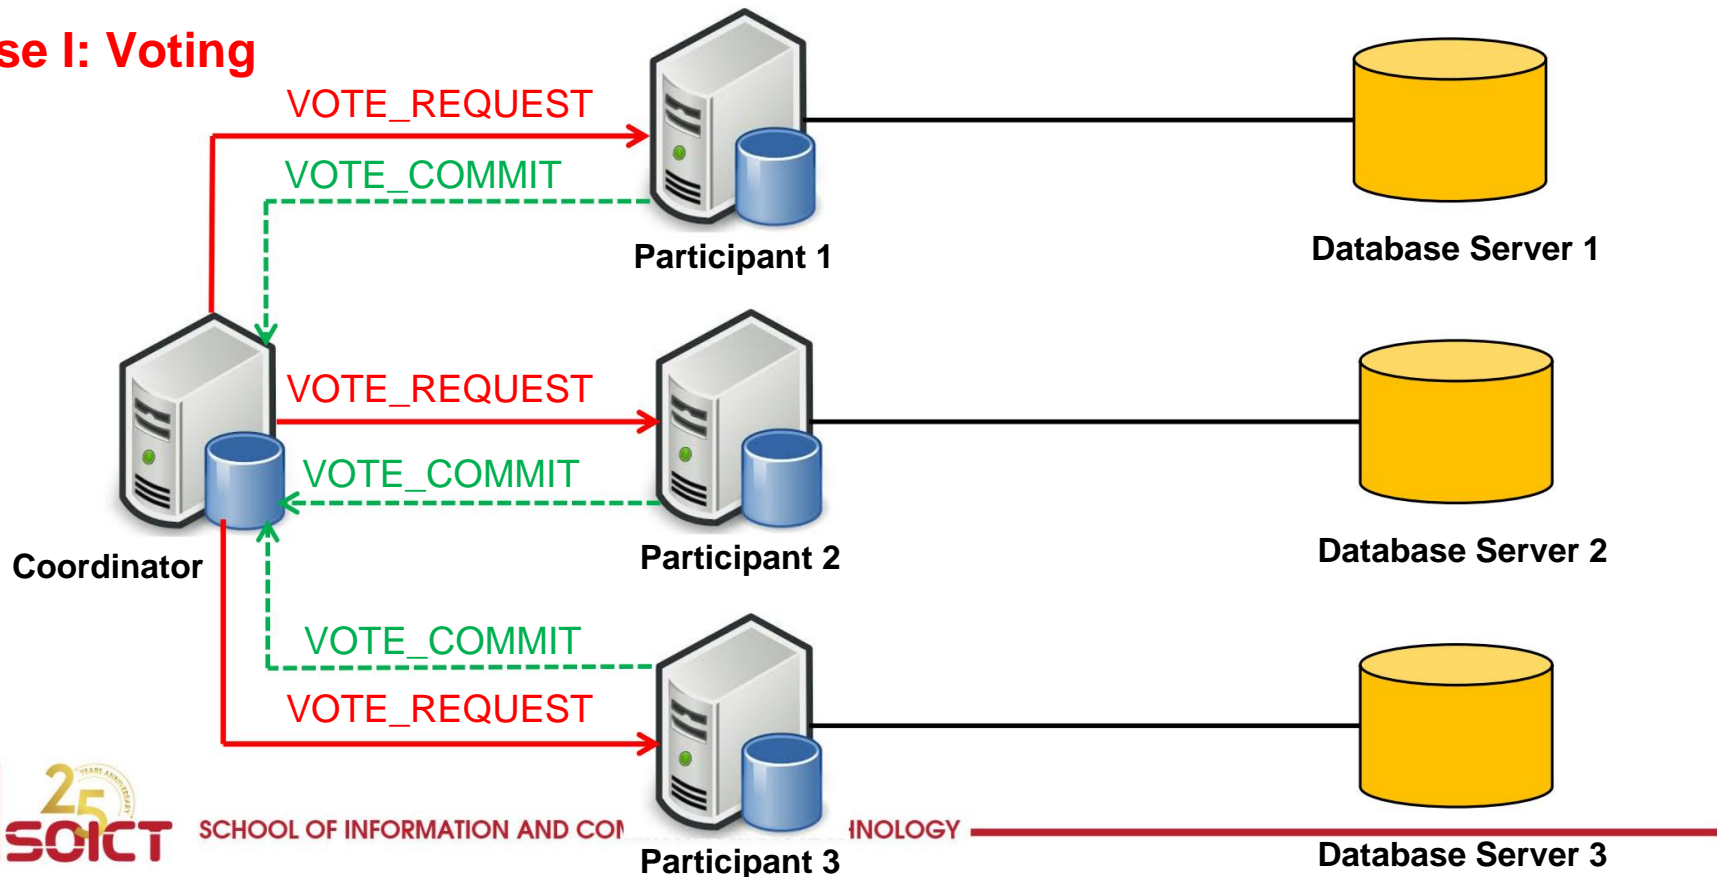
- An example:
  - In an e-commerce application, the bank database has been replicated across two servers
  - Maintaining consistency of replicated data is a challenge



# The Two-Phase Commit Protocol

- The two-phase commit protocol (2PC) can be used to ensure atomicity and consistency

## Phase I: Voting

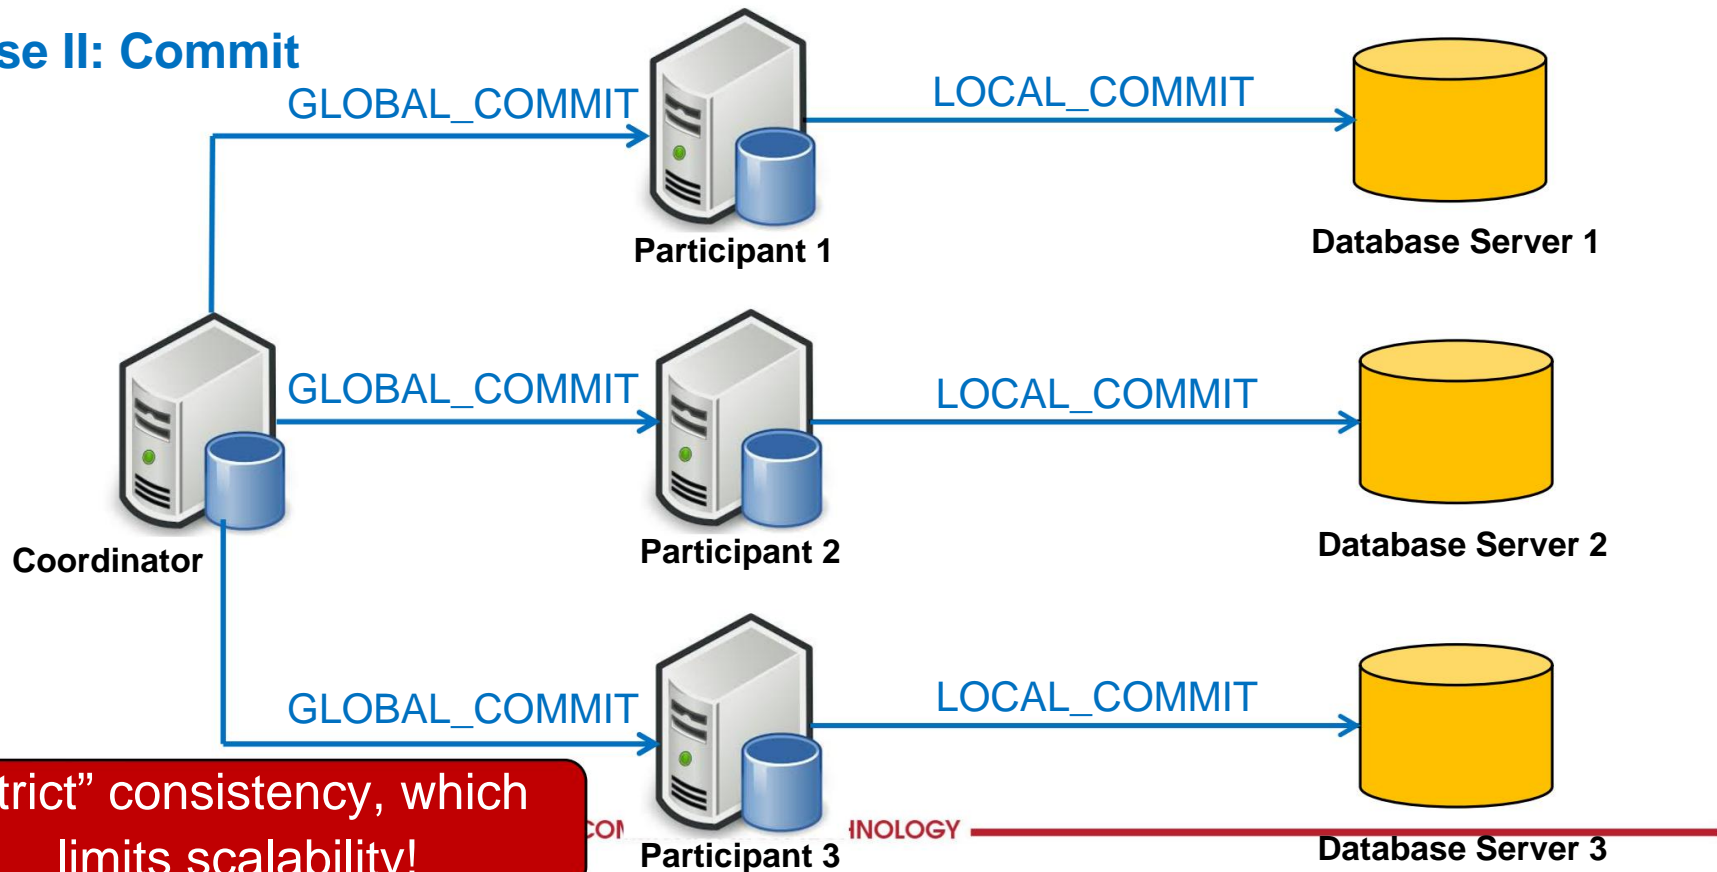




# The Two-Phase Commit Protocol

- The two-phase commit protocol (2PC) can be used to ensure atomicity and consistency

## Phase II: Commit



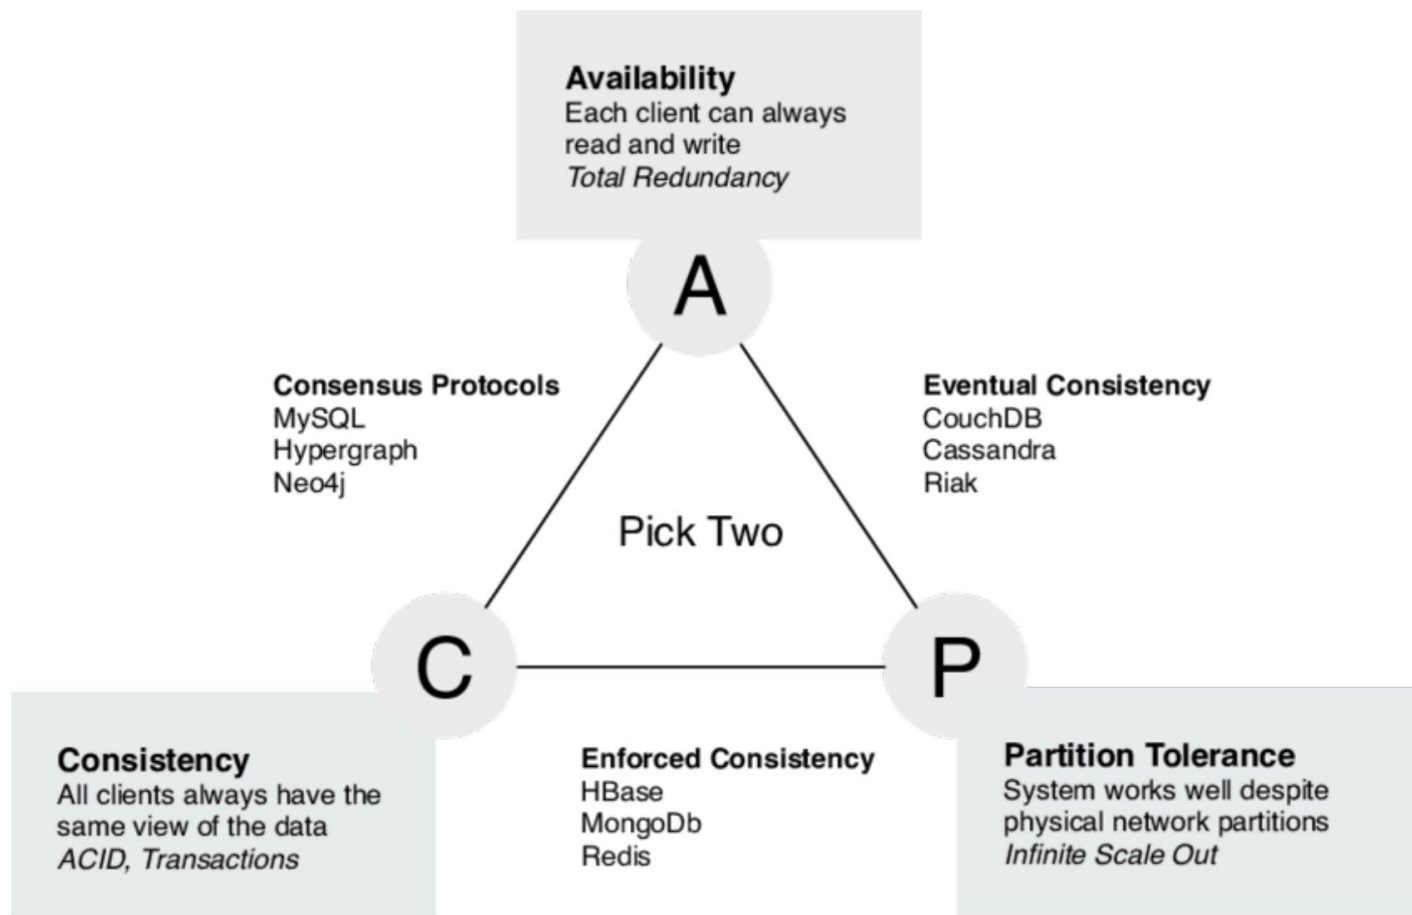


# The CAP Theorem

- The limitations of distributed databases can be described in the so called the CAP theorem
  - Consistency: every node always sees the same data at any given instance (ie, strict consistency)
  - Availability: the system continues to operate, even if nodes in a cluster crash, or some hardware or software parts are down due to upgrade
  - Partition Tolerance: the system continues to operate in the presence of network partitions


CAP theorem: any distributed database with shared data, can have at most two of the three desirable properties, C, A or P. **These are trade-offs involved in distributed system by Eric Brewer in PODC 2000.**

# CAP Theorem




# CAP Theorem: Proof

- A simple proof using two nodes:




# CAP Theorem: Proof

- A simple proof using two nodes:




# CAP Theorem: Proof

- A simple proof using two nodes:



# CAP Theorem: Proof

- A simple proof using two nodes:



A gets updated from B

# Scalability of relational databases

- The Relational Database is built on the principle of **ACID** (Atomicity, Consistency, Isolation, Durability)
- It implies that a truly distributed relational database should have **availability, consistency and partition tolerance.**
- Which unfortunately is **impossible** ...



# Large-Scale Databases

- When companies such as Google and Amazon were designing large-scale databases, 24/7 Availability was a key
  - A few minutes of downtime means lost revenue
- When horizontally scaling databases to 1000s of machines, the likelihood of a node or a network failure increases enormously •

Therefore, in order to have strong guarantees on Availability and Partition Tolerance, they had to sacrifice “strict” Consistency (implied by the CAP theorem)

# Trading-Off Consistency

- Maintaining consistency should balance between the strictness of consistency versus availability/scalability
  - Good-enough consistency depends on your application

# Trading-Off Consistency

- Maintaining consistency should balance between the strictness of consistency versus availability/scalability
  - Good-enough consistency depends on your application



# The BASE Properties

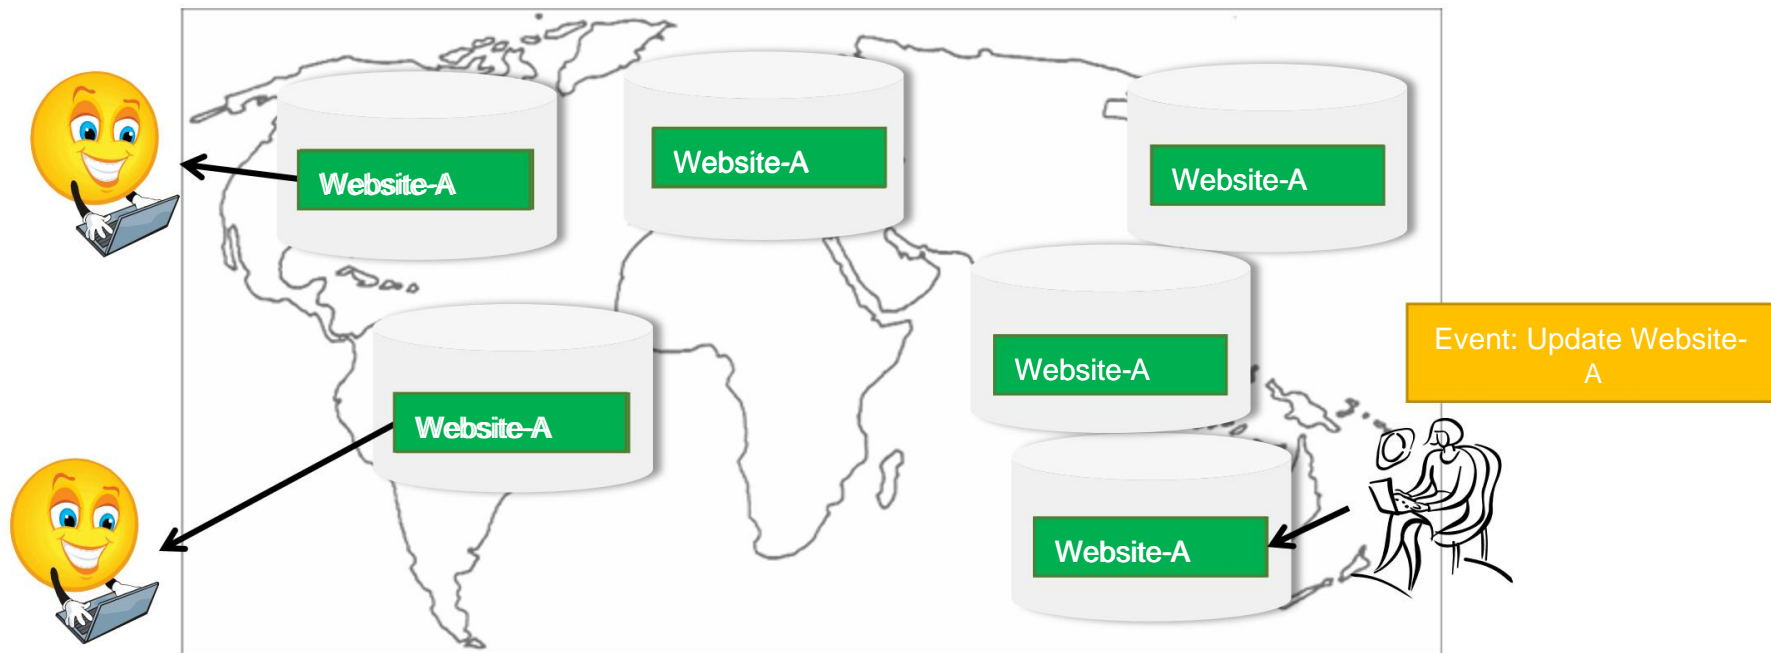
- The CAP theorem proves that it is impossible to guarantee strict Consistency and Availability while being able to tolerate network partitions
- This results in databases with relaxed ACID guarantees
- In particular, such databases apply the BASE properties:
  - Basically Available: the system guarantees Availability
  - Soft-State: the state of the system may change over time
  - Eventual Consistency: the system will eventually become consistent

# Eventual Consistency

- A database is termed as eventually Consistent if:
  - All replicas will gradually become consistent in the absence of new updates

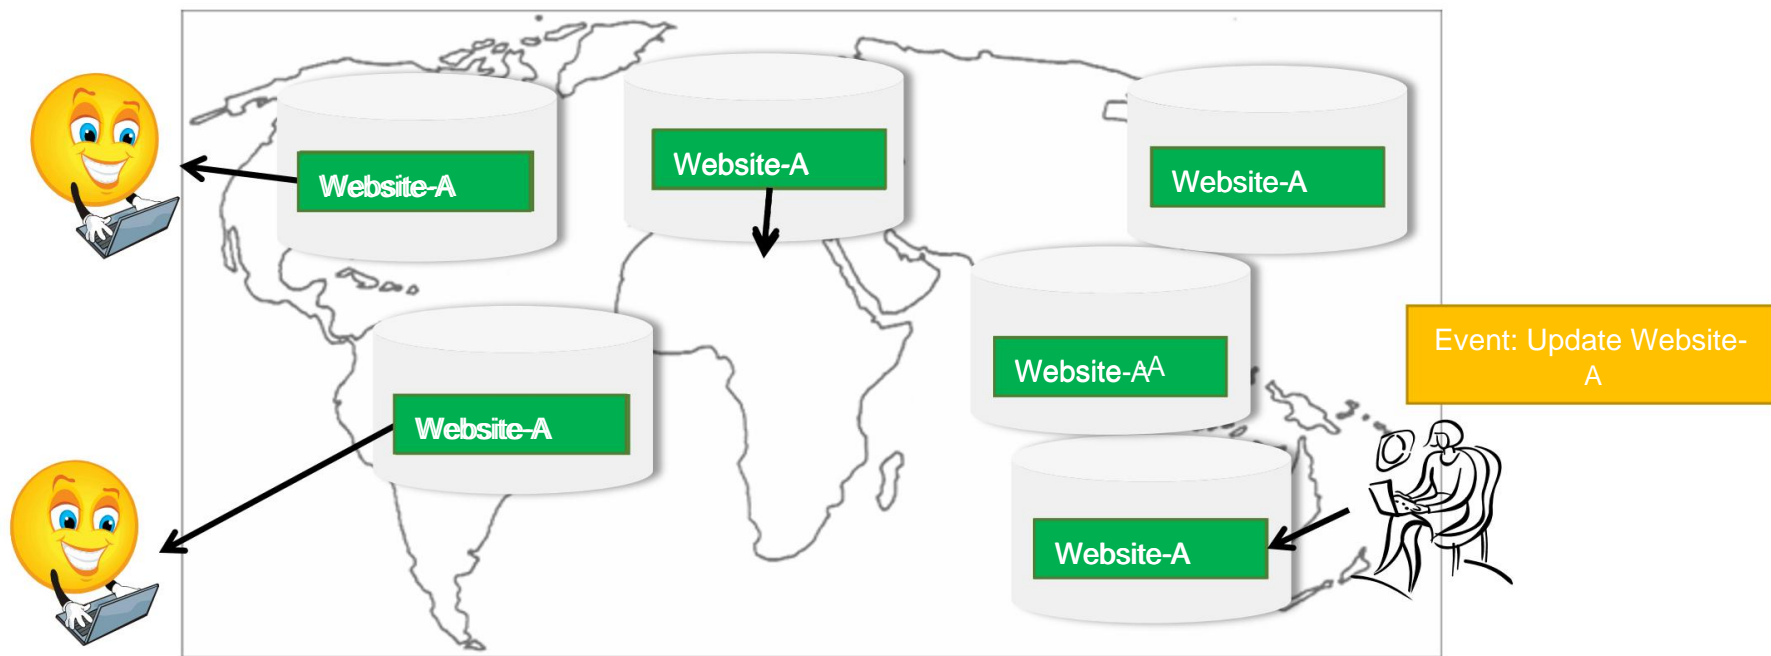
# Eventual Consistency

- A database is termed as eventually Consistent if:
  - All replicas will gradually become consistent in the absence of new updates



# Read-after-write consistency (eg. Amazon S3)

- But, what if the client accesses the data from different replicas?



Protocols like Read Your Own Writes (RYOW) can be applied!



# References

- Gilbert, Seth, and Nancy Lynch. "Perspectives on the CAP Theorem." *Computer* 45.2 (2012): 30-36.
- Kleppmann, Martin. "A Critique of the CAP Theorem." *arXiv preprint arXiv:1509.05393* (2015).
- Brewer, Eric. "CAP twelve years later: How the "rules" have changed." *Computer* 45.2 (2012): 23-29.
- Chandra, Deka Ganesh. "BASE analysis of NoSQL database." *Future Generation Computer Systems* 52 (2015): 13-21.



VIỆN CÔNG NGHỆ THÔNG TIN VÀ TRUYỀN THÔNG  
SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

Thank you  
for your  
attention!!!



[soict.hust.edu.vn/](http://soict.hust.edu.vn/)



[fb.com/groups/soict](https://fb.com/groups/soict)

